

Compression

- What
 - Reduce the amount of information (bits) needed to represent image
 - Video: 720 x 480 res, 30 fps, color
 - $720 \times 480 \times 20 \times 3 = 31,104,000$ bytes/sec
 - $30 \times 60 \times 120 = 216$ Gigabytes for a 2 hour movie
- Why
 - Transmission (send video over wireless channel)
 - Storage (fit a 2 hour movie on a DVD)

Compression Model

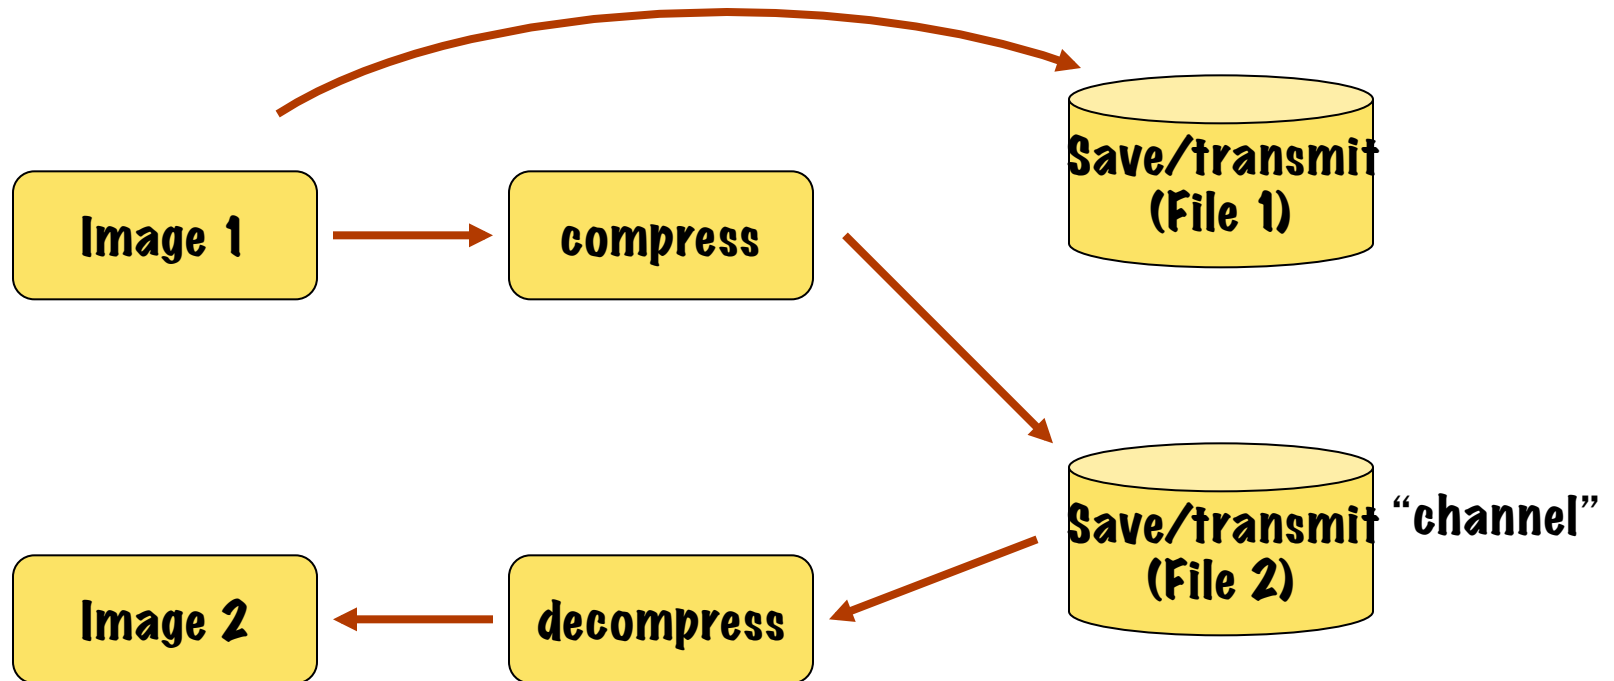


Image1 == Image2 -> "lossless" <- reduces redundant info


Image1 != Image2 -> "lossy" <- tries to reduce redundant & irrelevant info

Size(File1)/Size(File2) -> "compression ratio"

Redundancy

- Coding redundancy
 - More bits than necessary to create unique codes
- Spatiotemporal redundancy
 - Correlation between pixels
 - Patterns in image, motion in video
- Irrelevant information
 - Human visual system cannot distinguish more than a certain number of gray levels in a given image

1. Coding redundancy



$$p_r(r_k) = \frac{n_k}{MN}$$

$$L_{avg} = \sum_{k=0}^{L-1} l(r_k) p_r(r_k)$$

$l(r_k)$ Number of bits needed for level k

| r_k | $p_r(r_k)$ | Code 1 | $l_1(r_k)$ | Code 2 | $l_2(r_k)$ |
|--------------------------------------|------------|----------|------------|--------|------------|
| $r_{87} = 87$ | 0.25 | 01010111 | 8 | 01 | 2 |
| $r_{128} = 128$ | 0.47 | 10000000 | 8 | 1 | 1 |
| $r_{186} = 186$ | 0.25 | 11000100 | 8 | 000 | 3 |
| $r_{255} = 255$ | 0.03 | 11111111 | 8 | 001 | 3 |
| r_k for $k \neq 87, 128, 186, 255$ | 0 | — | 8 | — | 0 |

Code 1: $L_{avg} = 8$ bits

Code 2: $L_{avg} = 0.25 \cdot 2 + 0.47 \cdot 1 + 0.25 \cdot 3 + 0.03 \cdot 3 = 1.81$ bits

Compression ratio $C = 8/1.81 = 4.42$


Information Theory

- How much information does a random event E give us?
 - $I(E) = \log[1/P(E)] = -\log P(E)$
 - Base 2 log, unit of information is **bits**
 - If an event is sure to happen, i.e. $P(E)=1$, do I get much information when it happens?
 - *Event: The sun rose this morning. Info: The world didn't end?*
 - If an event is unlikely to happen, small $P(E)$, if it happens we get a lot of information.
 - *Event: It is raining in my bedroom. Info: my roof must be leaking.*

Information Theory

- Information content of a signal -> entropy

$$H = -\sum_{j=0}^{L-1} P(r_j) \log(P(r_j)) \quad H = -\int p(x) \log p(x) dx$$



- Entropy gives the lower bound on #bits need to unambiguously represent a sequence of symbols

Back to coding example

| r_k | $p_r(r_k)$ | Code 1 | $l_1(r_k)$ | Code 2 | $l_2(r_k)$ |
|--------------------------------------|------------|----------|------------|--------|------------|
| $r_{87} = 87$ | 0.25 | 01010111 | 8 | 01 | 2 |
| $r_{128} = 128$ | 0.47 | 10000000 | 8 | 1 | 1 |
| $r_{186} = 186$ | 0.25 | 11000100 | 8 | 000 | 3 |
| $r_{255} = 255$ | 0.03 | 11111111 | 8 | 001 | 3 |
| r_k for $k \neq 87, 128, 186, 255$ | 0 | — | 8 | — | 0 |

Code 1: $L_{\text{avg}} = 8$ bits / pixel

Code 2: $L_{\text{avg}} = 0.25 \cdot 2 + 0.47 \cdot 1 + 0.25 \cdot 3 + 0.03 \cdot 3 = 1.81$ bits / pixel

Compression ratio $C = 8 / 1.81 = 4.42$

What does information theory say? (all logs are base 2)

$$H = -[0.25 \log 0.25 + 0.47 \log 0.47 + 0.25 \log 0.25 + 0.03 \log 0.03]$$

H=1.6614 bits / pixel. This is a lower bound. Could be achieved theoretically if coding groups of symbols rather than one at a time

Huffman Coding

- Need a method for computing the code
- Optimal when coding one symbol at a time
 - Variable length code
- Source reduction step
 - Order probabilities
 - Combine lowest probability pair into compound symbol, repeat.

| Original source | | Source reduction | | | |
|-----------------|-------------|------------------|-----|-----|-----|
| Symbol | Probability | 1 | 2 | 3 | 4 |
| a_2 | 0.4 | 0.4 | 0.4 | 0.4 | 0.6 |
| a_6 | 0.3 | 0.3 | 0.3 | 0.3 | |
| a_1 | 0.1 | 0.1 | 0.2 | 0.3 | 0.4 |
| a_4 | 0.1 | 0.1 | | | |
| a_3 | 0.06 | 0.1 | 0.1 | | |
| a_5 | 0.04 | | | | |

Huffman Coding

- Code assignment step
 - Start at the end, assign 0 and 1 to the two compound symbols, work backwards

| Original source | | | Source reduction | | | |
|-----------------|-------------|-------|------------------|------|-----|-----|
| Symbol | Probability | Code | 1 | 2 | 3 | 4 |
| a_2 | 0.4 | 1 | 0.4 | 1 | 0.4 | 1 |
| a_6 | 0.3 | 00 | 0.3 | 00 | 0.3 | 00 |
| a_1 | 0.1 | 011 | 0.1 | 011 | 0.2 | 010 |
| a_4 | 0.1 | 0100 | 0.1 | 0100 | 0.1 | 011 |
| a_3 | 0.06 | 01010 | 0.1 | 0101 | 0.3 | 01 |
| a_5 | 0.04 | 01011 | | | 0.6 | 0 |
| | | | | | 0.4 | 1 |

$$L_{\text{avg}} = 0.4 * 1 + 0.3 * 2 + 0.1 * 3 + 0.1 * 4 + 0.06 * 5 + 0.04 * 5 = 2.2 \text{ bits/pixel}$$

$$H = [0.4 \log 0.4 + 0.3 \log 0.3 + 0.1 \log 0.1 + 0.1 * \log 0.01 + 0.06 \log 0.06 + 0.04 \log 0.04] = 2.1435 \text{ bits/pixel}$$

Huffman coding

- After the code is created, it serves as a look-up table for coding and lossless decoding
- Block code: each source symbol is mapped into a fixed sequence of code symbols
- Instantaneous: Each code word can be decoded without reference to previous symbols in the sequence.
- Unique: Any sequence of symbols can be decoded in only one way.

Coding example

| Original source | | | Source reduction | | | | | | | |
|-----------------|-------------|-------|------------------|------|-----|-----|-----|----|-----|---|
| Symbol | Probability | Code | 1 | | 2 | | 3 | | 4 | |
| a_2 | 0.4 | 1 | 0.4 | 1 | 0.4 | 1 | 0.4 | 1 | 0.6 | 0 |
| a_6 | 0.3 | 00 | 0.3 | 00 | 0.3 | 00 | 0.3 | 00 | 0.4 | 1 |
| a_1 | 0.1 | 011 | 0.1 | 011 | 0.2 | 010 | 0.3 | 01 | | |
| a_4 | 0.1 | 0100 | 0.1 | 0100 | 0.1 | 011 | | | | |
| a_3 | 0.06 | 01010 | 0.1 | 0101 | | | | | | |
| a_5 | 0.04 | 01011 | | | | | | | | |

- Generate the code for the sequence $a_3 a_1 a_2 a_2 a_6$
 - a_3 a_1 a_2 a_2 a_6
 - 01010 011 1 1 00
- 010100111100


Decoding example

| Original source | | | Source reduction | | | | | | | |
|-----------------|-------------|-------|------------------|------|-----|-----|-----|----|-----|---|
| Symbol | Probability | Code | 1 | | 2 | | 3 | | 4 | |
| a_2 | 0.4 | 1 | 0.4 | 1 | 0.4 | 1 | 0.4 | 1 | 0.6 | 0 |
| a_6 | 0.3 | 00 | 0.3 | 00 | 0.3 | 00 | 0.3 | 00 | 0.4 | 1 |
| a_1 | 0.1 | 011 | 0.1 | 011 | 0.2 | 010 | 0.3 | 01 | | |
| a_4 | 0.1 | 0100 | 0.1 | 0100 | 0.1 | 011 | | | | |
| a_3 | 0.06 | 01010 | 0.1 | 0101 | | | | | | |
| a_5 | 0.04 | 01011 | | | | | | | | |

- Decode 010100111100.
 - Find first valid code word 010100111100 -> a_3
 - Find next valid code word 010100111100 -> $a_3 a_1$
 - 01010 011 1 1 00 -> $a_3 a_1 a_2 a_2 a_6$
- $a_3 a_1 a_2 a_2 a_6$

2. Spatial redundancy

- Images have homogeneous regions
- Run-length encoding
 - Row-major order
 - Encode value of “run” and it’ s length
 - Can combine with symbol encoder
- Good for images with few, discrete color values. For instance, binary images
- Issues
 - How homogeneous is the data?
 - Is there enough continuity in rows?



BMP, JPEG, MPEG

BMP file format

- Uses a form of RLE. Two modes
 - Encoded: 2 byte representation: first byte number of pixels, second byte value/color index
 - Absolute: First byte 0. Second byte has special meaning

| Second Byte Value | Condition |
|-------------------|-----------------------------|
| 0 | End of line |
| 1 | End of image |
| 2 | Move to a new position |
| 3–255 | Specify pixels individually |

→ Next 2 bytes position

→ This many
uncompressed
pixels follow

Fixed Length Codes

- Dictionary with strategy to capture special structure of data
- Example: LZW (Lempel-Ziv-Welch)
 - Start with basic dictionary (e.g. grey levels)
 - As new sequences of symbols are encountered add them to dictionary
 - Hope: encode frequently occurring sequences of symbols
 - Greedy
 - Can decompress w/out table

LZW coding

| | | | |
|-----------|-----------|------------|------------|
| 39 | 39 | 126 | 126 |
| 39 | 39 | 126 | 126 |
| 39 | 39 | 126 | 126 |
| 39 | 39 | 126 | 126 |

4 x 4 image
512-word dictionary
0-255 stores intensities
256-511 initially unused
Process left-to-right, top-to-
bottom fashion

| Dictionary Location | Entry |
|----------------------------|--------------|
| 0 | 0 |
| 1 | 1 |
| ⋮ | ⋮ |
| 255 | 255 |
| 256 | — |
| ⋮ | ⋮ |
| 511 | — |

| | | | |
|-----------|-----------|------------|------------|
| 39 | 39 | 126 | 126 |
| 39 | 39 | 126 | 126 |
| 39 | 39 | 126 | 126 |
| 39 | 39 | 126 | 126 |



4 x 4 image
512-word dictionary
0-255 stores intensities
256-511 initially unused
Process left-to-right, top-
to-bottom fashion

| Currently Recognized Sequence | Pixel Being Processed | Encoded Output | Dictionary Location (Code Word) | Dictionary Entry |
|-------------------------------------|--------------------------|-------------------|---------------------------------------|------------------|
| | 39 | | | |
| 39 | 39 | 39 | 256 | 39-39 |
| 39 | 126 | 39 | 257 | 39-126 |
| 126 | 126 | 126 | 258 | 126-126 |
| 126 | 39 | 126 | 259 | 126-39 |
| 39 | 39 | | | |
| 39-39 | 126 | 256 | 260 | 39-39-126 |
| 126 | 126 | | | |
| 126-126 | 39 | 258 | 261 | 126-126-39 |
| 39 | 39 | | | |
| 39-39 | 126 | | | |
| 39-39-126 | 126 | 260 | 262 | 39-39-126-126 |
| 126 | 39 | | | |
| 126-39 | 39 | 259 | 263 | 126-39-39 |
| 39 | 126 | | | |
| 39-126 | 126 | 257 | 264 | 39-126-126 |
| 126 | | 126 | | |

GIF , TIFF , PDF

3. Irrelevant information/ Quantization


- Eliminate symbols that are too small or not important
- Find a small set of approximating symbols (less entropy)
 - Grey level or “vector quantization”
 - Find values that minimize error
 - Related to “clustering” in pattern recognition




- Quantization corresponds to a transformation $Q(f)$

Code:

4 levels



8 levels



Could also use Huffman coding to generate a more optimal code using the histogram of the image

Fidelity criteria

$$SNR_{ms} = \frac{\sum_{x=0}^{M-1} \sum_{y=0}^{N-1} \hat{f}(x, y)^2}{\sum_{x=0}^{M-1} \sum_{y=0}^{N-1} [\hat{f}(x, y) - f(x, y)]^2}$$

Approximation
For instance
quantized image

Input

Can we choose the quantization transform $Q(f)$ to minimize SNR?

Lloyd-Max algorithm

- Bin boundaries $\min x = L_1 < L_2 < \dots < L_{n+1} = \max x$
- Replacement values p_1, p_2, \dots, p_n
- $q(x) = p_j$ when $L_j \leq x < L_{j+1}$
- Choose L and p to minimize

$$E(L, p) = \sum_{i=1}^m (x_i - q(x_i))^2 = \sum_{j=1}^n \sum_{x_i \in [L_j, L_{j+1})} (x_i - p_j)^2$$

Lloyd-Max algorithm

$$E(L, p) = \sum_{j=1}^n \sum_{x_i \in [L_j, L_{j+1})} (x_i - p_j)^2$$

- Derivative with respect to p

$$\frac{\partial E}{\partial p_j} = \sum_{x_i \in [L_j, L_{j+1})} 2(x_i - p_j) = 0$$

$$p_j = \frac{\sum_{x_i \in [L_j, L_{j+1})} x_i}{\# \{i | x_i \in [L_j, L_{j+1})\}}$$

Size of the set

Update rule for p_j

Average intensity of pixels
with intensities in this
interval

Lloyd-Max algorithm

$$E(L, p) = \sum_{j=1}^n \sum_{x_i \in [L_j, L_{j+1})} (x_i - p_j)^2$$

- Update rule for L
 - $L_j = 0.5 * (p_{j-1} + p_j)$. Why?
 - If you decrease L_j , you will be assigning some x_i to p_j even though it is closer to p_{j-1}
 - If you increase L_j , you will be assigning some x_i to p_{j-1} even though it is closer to p_j

Lloyd-Max algorithm

1. Start from uniform quantization bin boundaries

2. Update p using

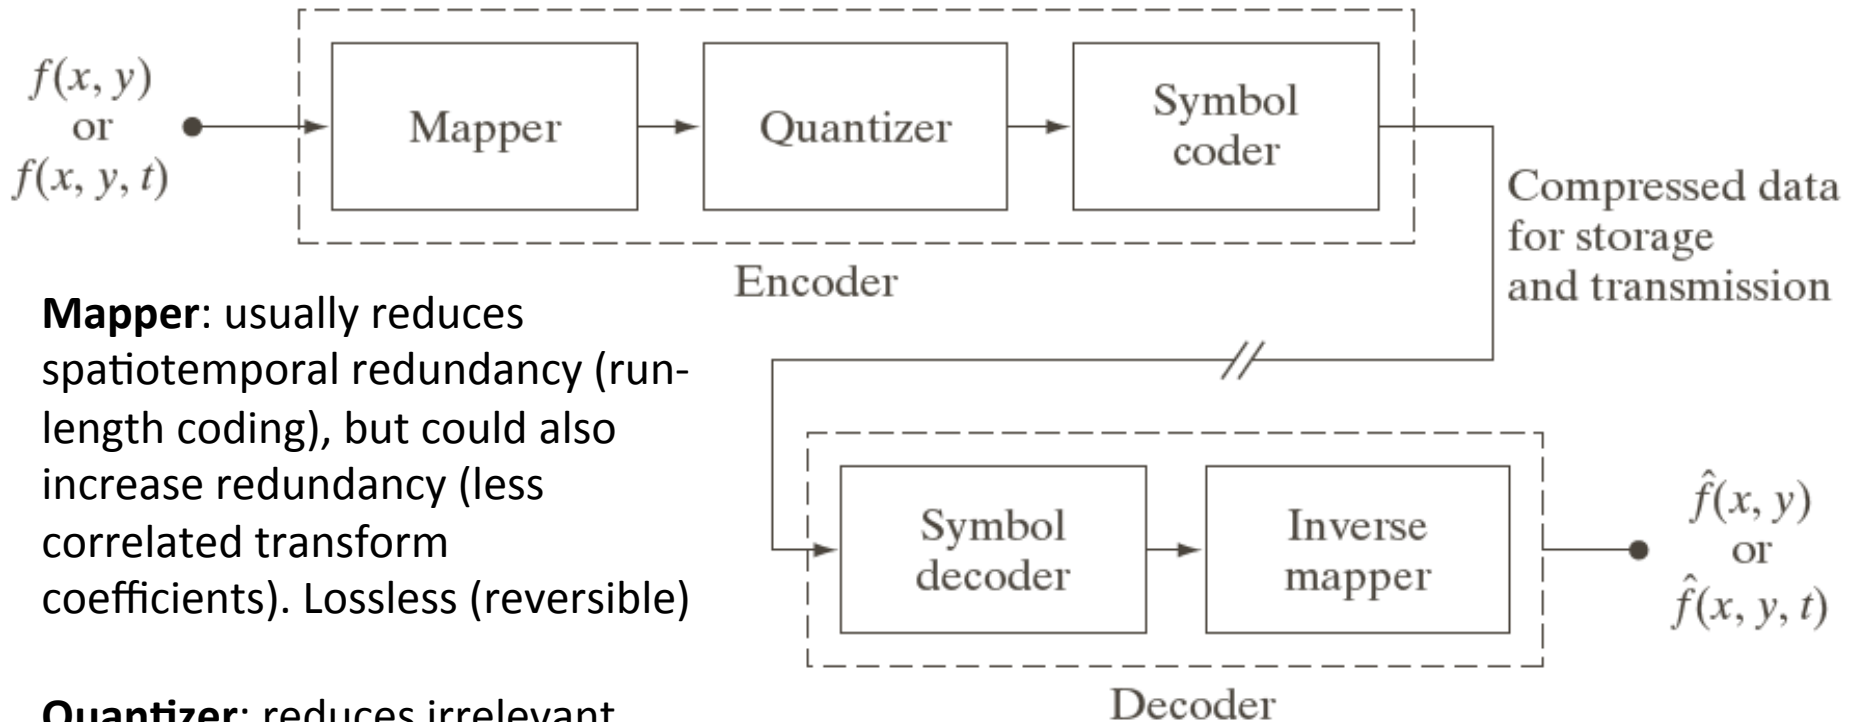
for j = 1 to n

$$p_j = \frac{\sum_{x_i \in [L_j, L_{j+1})} x_i}{\# \{i | x_i \in [L_j, L_{j+1})\}}$$

3. Update L using $L_j = 0.5 * (p_{j-1} + p_j)$ for j=2 to n. Notice lower and upper limit always set to min and max of image.

4. Go to step 2 until convergence

Image Compression Model



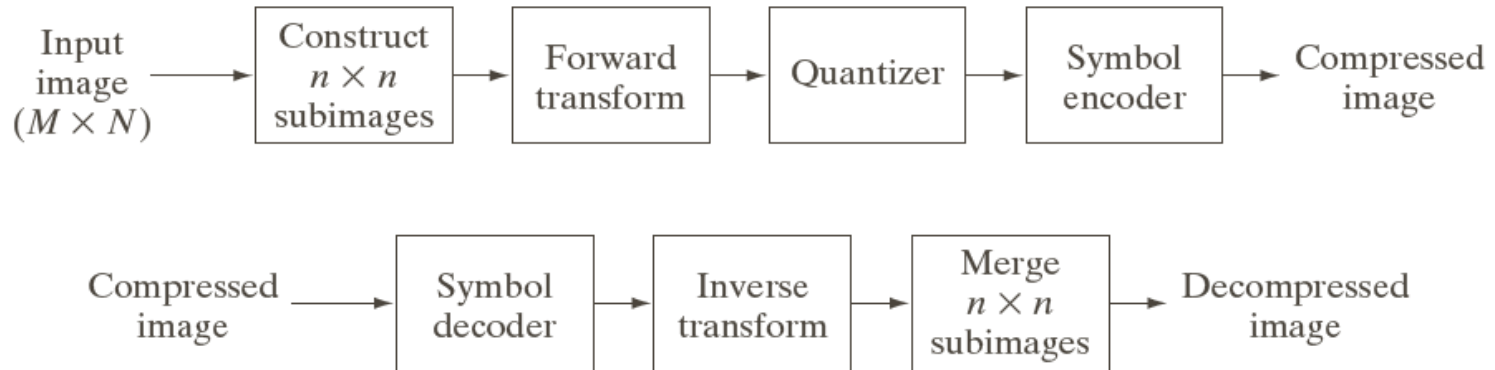
Mapper: usually reduces spatiotemporal redundancy (run-length coding), but could also increase redundancy (less correlated transform coefficients). Lossless (reversible)

Quantizer: reduces irrelevant information according to a fidelity criterion. Lossy

Symbol coder: For instance Huffman coding. Lossless.

Decoder: Undoes these in reverse order. Quantization can't be undone

Block transform coding



- **Forward transform:** decorrelate pixels + pack as much info as possible into fewest transform coefficients.
- **Quantization:** eliminate or coarsely quantize least informative transform coefficients
- **Symbol encoder:** Variable length coding for quantized coefficients
- All steps can be adaptive to subimages or fixed globally

Transforms

$$T(u, v) = \sum_{x=0}^{n-1} \sum_{y=0}^{n-1} g(x, y) r(x, y, u, v)$$

Forward transform

$$g(x, y) = \sum_{u=0}^{n-1} \sum_{v=0}^{n-1} T(u, v) s(x, y, u, v)$$

Inverse transform

Basis functions


Fourier basis:

$$r(x, y, u, v) = e^{-j2\pi(ux+vy)/n}$$

$$s(x, y, u, v) = (1/n^2) e^{j2\pi(ux+vy)/n}$$

Walsh Hadamard transform (WHT)

- A computationally very simple method. Basis functions consists of only +1 and -1
- Math for how they are produced in the textbook
- Notice frequency pattern
- Basis functions r and s are identical
- Example on the right is for $n=4$




Discrete Cosine Transform (DCT)

$$r(x, y, u, v) = s(x, y, u, v) = \alpha(u)\alpha(v) \cos\left[\frac{(2x+1)u\pi}{2n}\right] \cos\left[\frac{(2y+1)v\pi}{2n}\right]$$

$$\alpha(u) = \begin{cases} \sqrt{\frac{1}{n}} & u = 0 \\ \sqrt{\frac{2}{n}} & u = 1, 2, \dots, n-1 \end{cases}$$

Real valued basis functions
Used in JPEG



Fourier




Walsh Hadamard

Discrete Cosine

Reconstruction



Error




8 x 8 blocks: 64 coefficients, in each block 32 coefficients with smallest magnitude were truncated. Blocks then reconstructed from remaining 32 coefficients.

RMS Errors: Fourier 2.32, WHT 1.78, DCT 1.13

Subimage size selection

- n is usually a power of 2
- Larger n increases computational complexity (transforms don't scale linearly with number of pixels in subimage)



25% coefficients retained

How to truncate

| | | | | | | | |
|---|---|---|---|---|---|---|---|
| 1 | 1 | 1 | 1 | 1 | 0 | 0 | 0 |
| 1 | 1 | 1 | 1 | 0 | 0 | 0 | 0 |
| 1 | 1 | 1 | 0 | 0 | 0 | 0 | 0 |
| 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

| | | | | | | | |
|---|---|---|---|---|---|---|---|
| 1 | 1 | 0 | 1 | 1 | 0 | 0 | 0 |
| 1 | 1 | 1 | 1 | 0 | 0 | 0 | 0 |
| 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$$g(x, y) = \sum_{u=0}^{n-1} \sum_{v=0}^{n-1} T(u, v) X(u, v) s(x, y, u, v)$$

Top: zonal mask $X(u, v)$. Fixed strategy

Bottom: threshold mask. Depends on subimage.

Have to also know which elements we are keeping

How to quantize

- Zonal strategy for how many bits used per coefficient for quantization

| | | | | | | | |
|---|---|---|---|---|---|---|---|
| 8 | 7 | 6 | 4 | 3 | 2 | 1 | 0 |
| 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 6 | 5 | 4 | 3 | 3 | 1 | 1 | 0 |
| 4 | 4 | 3 | 3 | 2 | 1 | 0 | 0 |
| 3 | 3 | 3 | 2 | 1 | 1 | 0 | 0 |
| 2 | 2 | 1 | 1 | 1 | 0 | 0 | 0 |
| 1 | 1 | 1 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Block Transform Coding: JPEG

- International standard (ISO)
- Baseline algorithm with extensions
- Transform: discrete cosine transform (DCT)

- Encodes freq. Info w/out complex #s
- FT of larger, mirrored signal
- Does not have other nice prop. of FT

$$F_u = \alpha(u) \sum_{i=0}^{N-1} f_i \cos \left[\frac{(2i+1)u\pi}{2N} \right]$$

$$F_i = \sum_{u=0}^{N-1} \alpha(u) F_u \cos \left[\frac{(2i+1)u\pi}{2N} \right]$$

$$\alpha(u) = \begin{cases} \sqrt{\frac{1}{N}} & u = 0 \\ \sqrt{\frac{2}{N}} & u \neq 0 \end{cases}$$

JPEG Algorithm

- Integer grey-level image broken into 8x8 sub blocks
- Set middle (mean?) grey level to zero (subtract middle)
- DCT of sub blocks (11 bit precision) $\rightarrow T(u,v)$
- Rescale frequency components by $Z(u,v)$ and round

Rescaling

$$\hat{T}(u, v) = \text{round} \left(\frac{T(u, v)}{Z(u, v)} \right)$$

- Different scaling matrices possible, but recommended is:

$$Z(u, v) = \begin{bmatrix} 16 & 11 & 10 & 16 & 24 & 40 & 51 & 61 \\ 12 & 12 & 14 & 19 & 26 & 58 & 60 & 55 \\ 14 & 13 & 16 & 24 & 40 & 57 & 69 & 56 \\ 14 & 17 & 22 & 29 & 51 & 87 & 80 & 62 \\ 18 & 22 & 37 & 56 & 68 & 109 & 103 & 77 \\ 24 & 35 & 55 & 64 & 81 & 104 & 113 & 92 \\ 49 & 64 & 78 & 87 & 103 & 121 & 120 & 101 \\ 72 & 92 & 95 & 98 & 112 & 100 & 103 & 99 \end{bmatrix}$$

Reordering

- DCT entries reordered in zig-zag fashion to increase coherency (produce blocks of zeros)


| | | | | | | | |
|----|----|----|----|----|----|----|----|
| 0 | 1 | 5 | 6 | 14 | 15 | 27 | 28 |
| 2 | 4 | 7 | 13 | 16 | 26 | 29 | 42 |
| 3 | 8 | 12 | 17 | 25 | 30 | 41 | 43 |
| 9 | 11 | 18 | 24 | 31 | 40 | 44 | 53 |
| 10 | 19 | 23 | 32 | 39 | 45 | 52 | 54 |
| 20 | 22 | 33 | 38 | 46 | 51 | 55 | 60 |
| 21 | 34 | 37 | 47 | 50 | 56 | 59 | 61 |
| 35 | 36 | 48 | 49 | 57 | 58 | 62 | 63 |

Coding

- Each sub-block is coded as a difference from previous sub-block
- Zeros are run-length encoded and nonzero elements are Huffman coded
 - Modified HC to allow for zeros

JPEG Example

Compression Ratio ~10:1



Loss of high frequencies

Ringing

Block artifacts